

TO COMPLETE YOUR REGISTRATION, PLEASE TELL US
WHETHER OR NOT THIS IMAGE CONTAINS A STOP SIGN:



ANSWER QUICKLY—OUR SELF-DRIVING
CAR IS ALMOST AT THE INTERSECTION.

SO MUCH OF "AI" IS JUST FIGURING OUT WAYS
TO OFFLOAD WORK ONTO RANDOM STRANGERS.

Convolutional Neural Networks

Logistics

- Quiz 10 due this Friday (Apr. 25)
- Last assignment – assignment 5 due next Monday (Apr. 28)
- The last quiz will be released in a moment, it's going to be free response questions on today's lecture. I will give you some time to finish it here.
- Last class “online asynchronous”

Logistics

- Course evaluation: fill out your course evaluation, **upload a screenshot showing you have completed it for CSC532 C**; then you earn 1 bonus point; if more than 90% of students finish this, then everyone gets 1 additional point. So in total, 2 bonus points available.
- <https://uiscoursesurvey.apps.uis.edu>
- Course evaluation bonus assignments has been available on Canvas, and it's due **23:45, May 3**.
- <https://uispringfield.instructure.com/courses/15935/assignments/275198>

Logistics

- Last course survey (1 bonus point) will be released online around May. 5, and due May 9
- Participation scores (attendance and Ed Discussion) will be released around May 9, right after the end of all the participation activities.
- Final Project will be due on 23:59, May 9. I need time to review your projects and submit final grades on time. Thus, please submit your project on time.
- Late final projects will get **a penalty of 15% per day up to three days after the due date.** After that, you can no longer submit the project.

Image Classification

Recognize which digit it is!


$$Y = f(\text{img})$$

$$P(Y = 0 \mid \text{img}) \quad \dots \quad P(Y = 9 \mid \text{img})$$


Please note that the slides for CNNs are adapted from

<https://www.youtube.com/watch?v=FrKWiRv254g>

Image Classification



Neural Network -- Inference



Find the Difference between Prediction and Truth

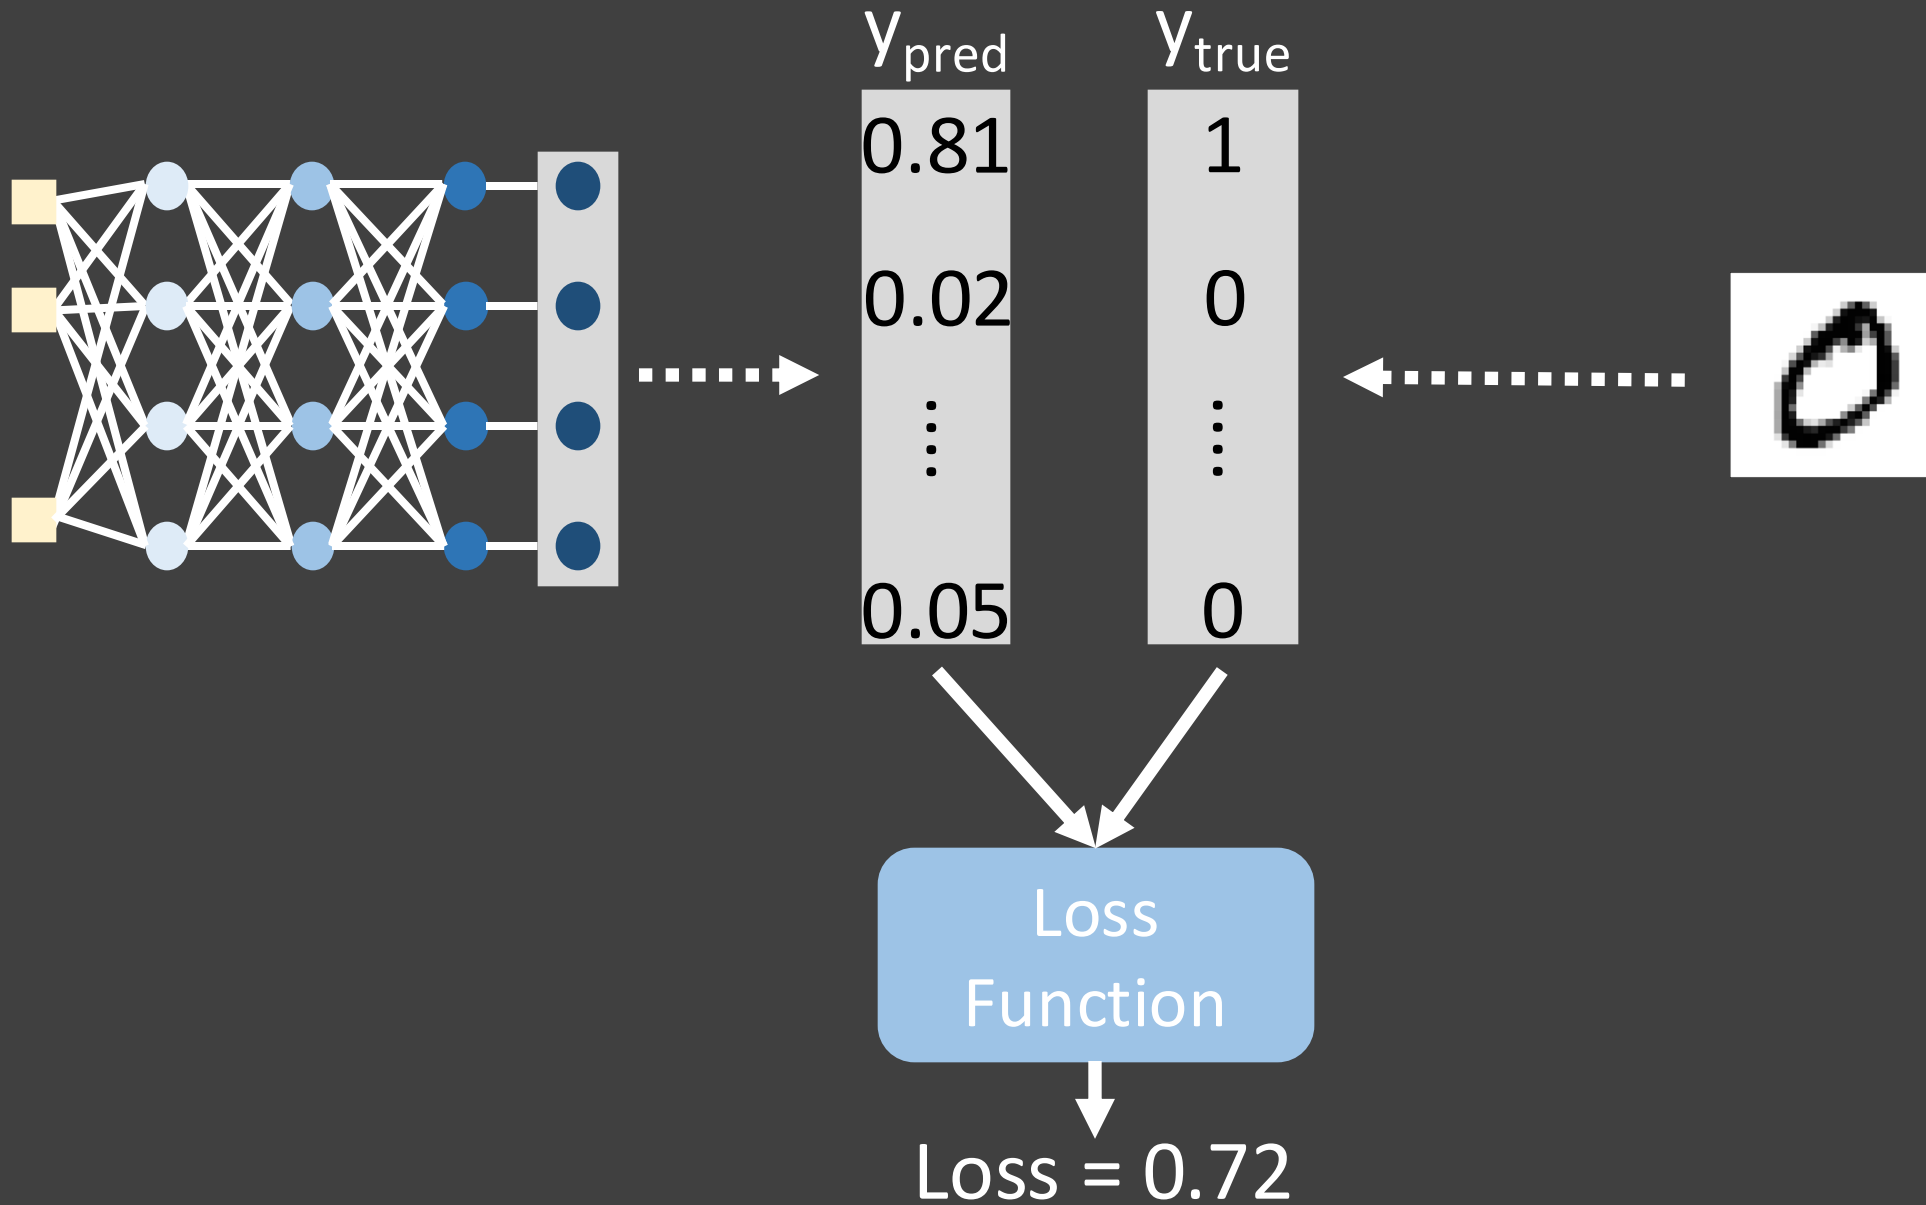



Image Classification



Input nodes = $6000 * 2000$
Hidden nodes per layer = 100
Too many parameters!!!

CNN saves the world

- Some patterns are much smaller than the whole image

A neuron does not have to see the whole image to discover the pattern.

Connecting to small region with less parameters

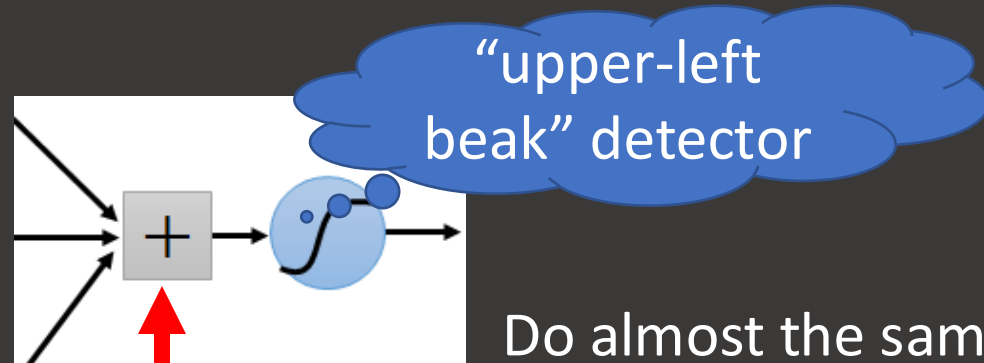


Slide adopted from

<https://www.youtube.com/watch?v=FrKWiRv254g>

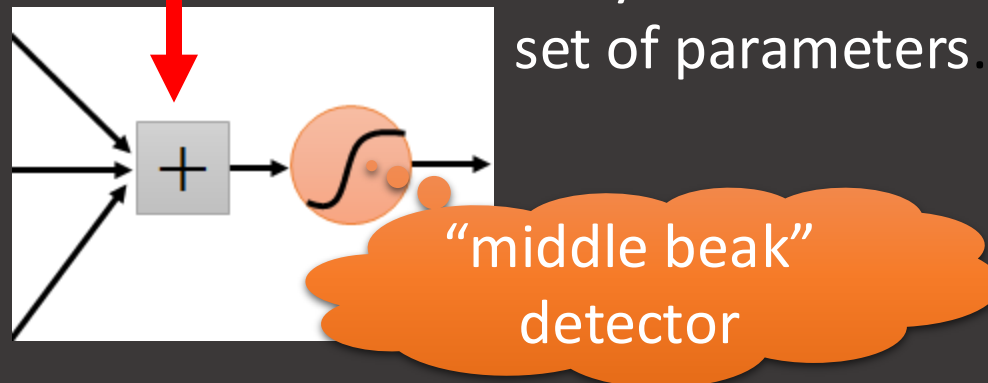
CNN saves the world

- The same patterns appear in different regions.



Do almost the same thing

They can use the same set of parameters.



CNN saves the world

- Subsampling the pixels will not change the object

bird



subsampling

bird



We can subsample the pixels to make image smaller

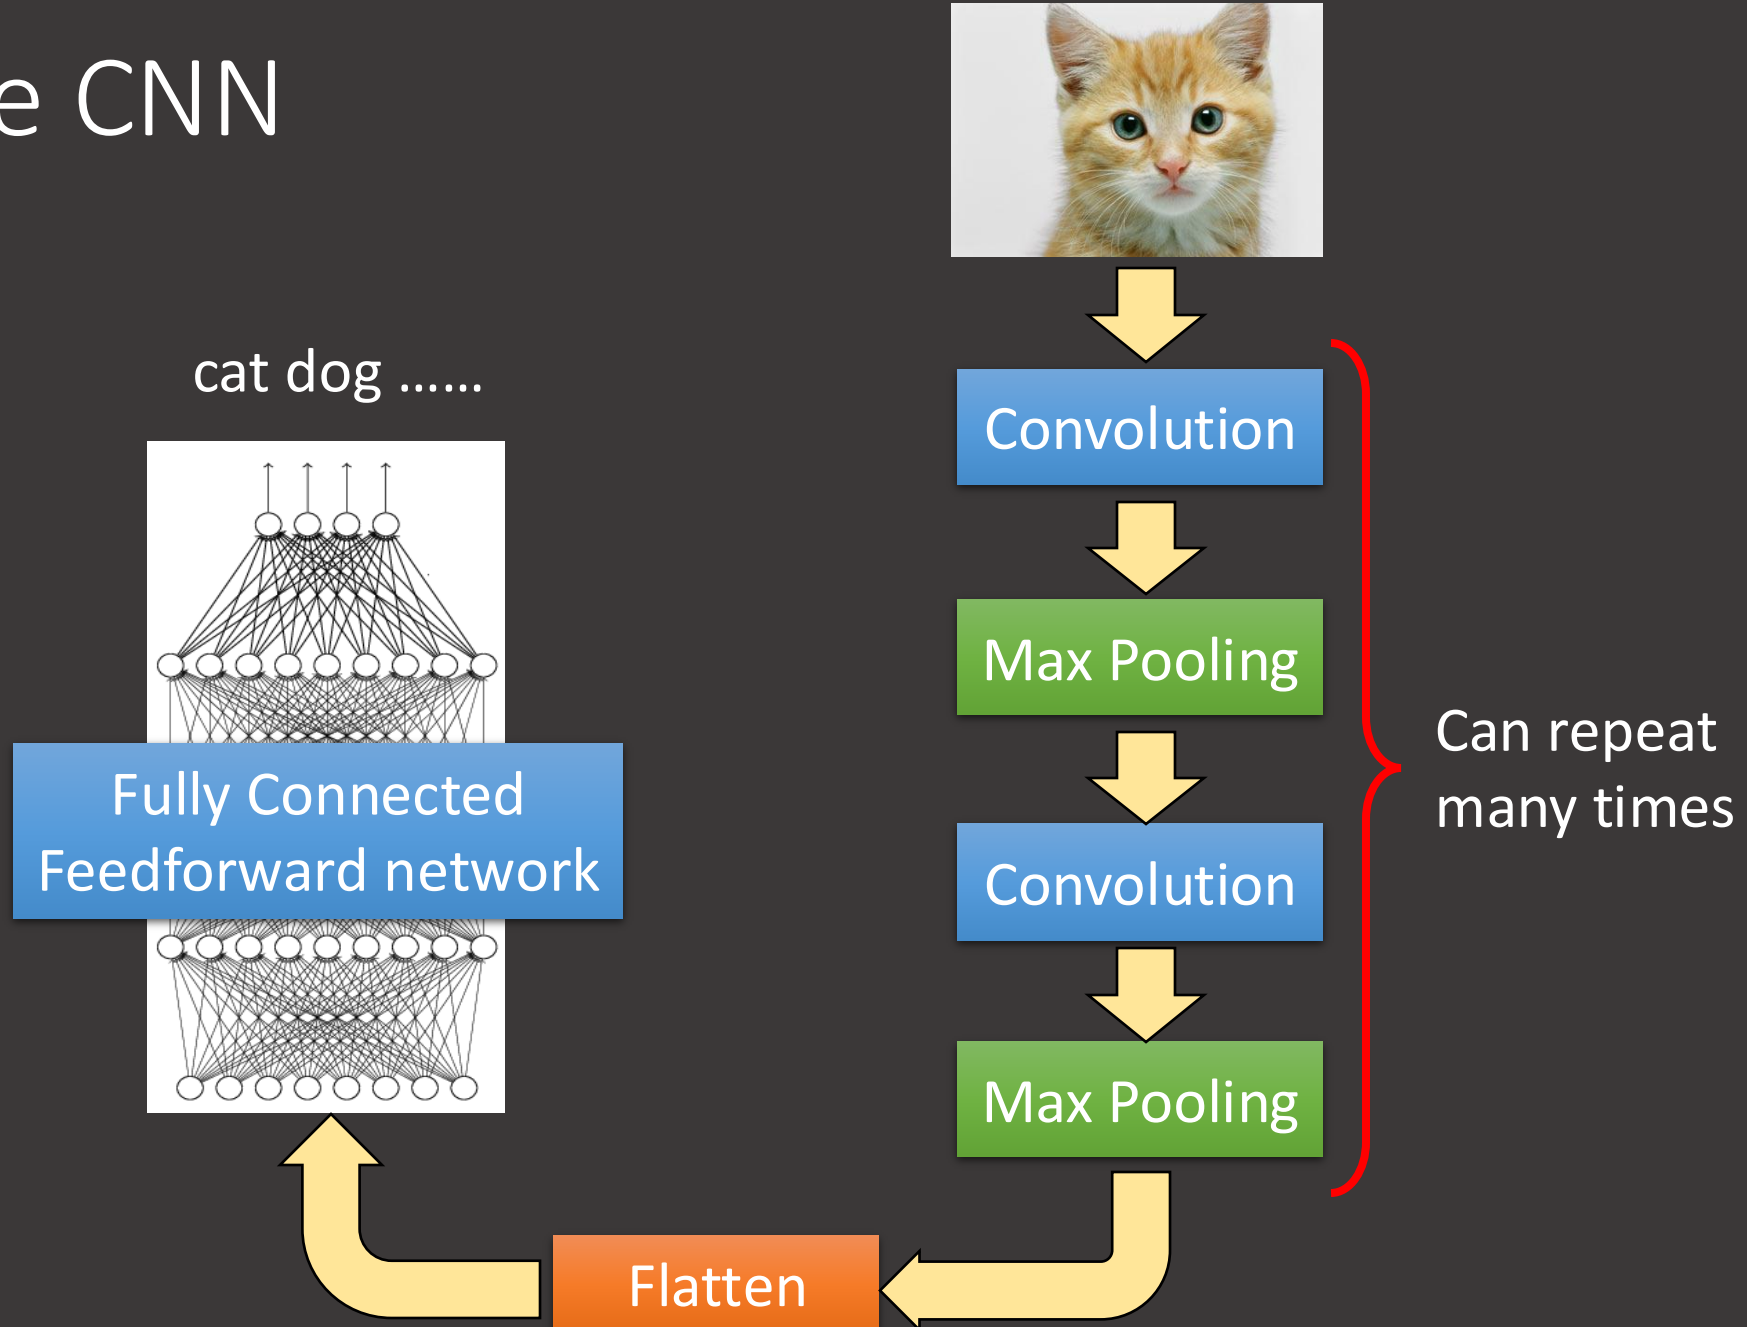


Less parameters for the network to process the image

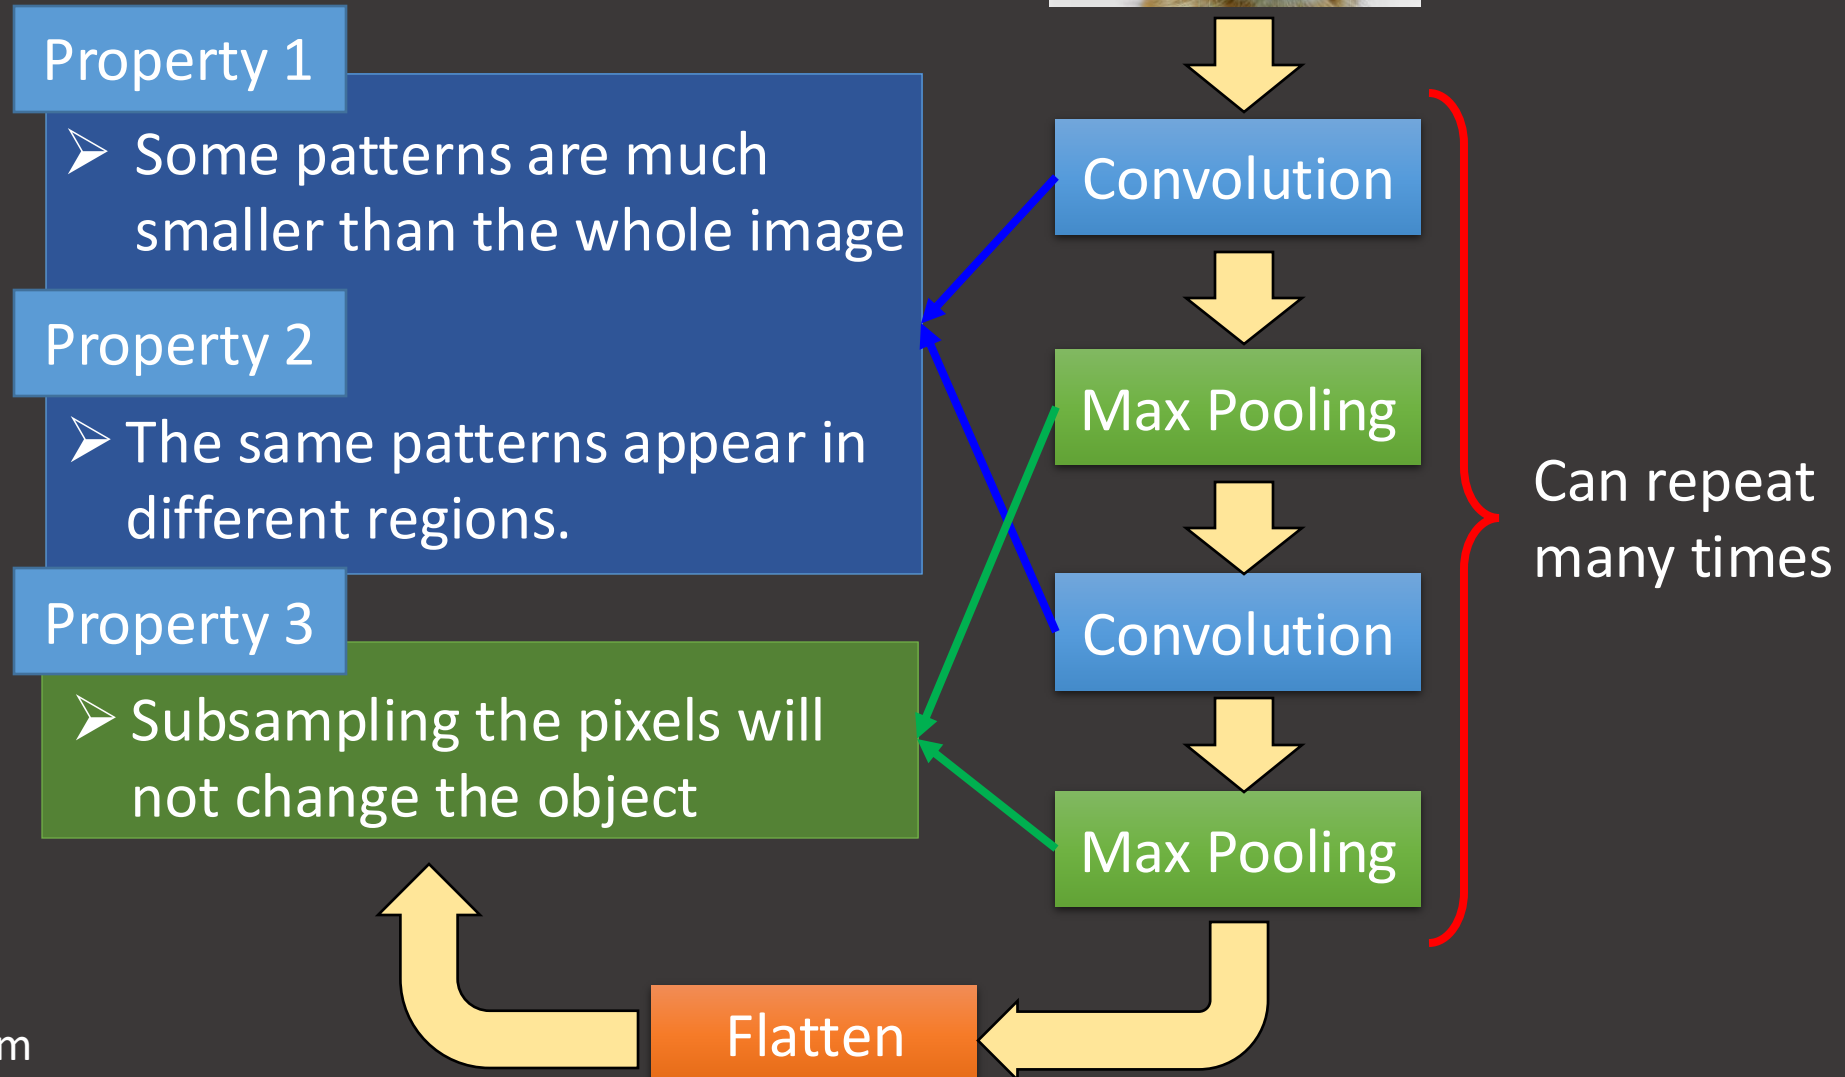
Slide adopted from

<https://www.youtube.com/watch?v=FrKWiRv254g>

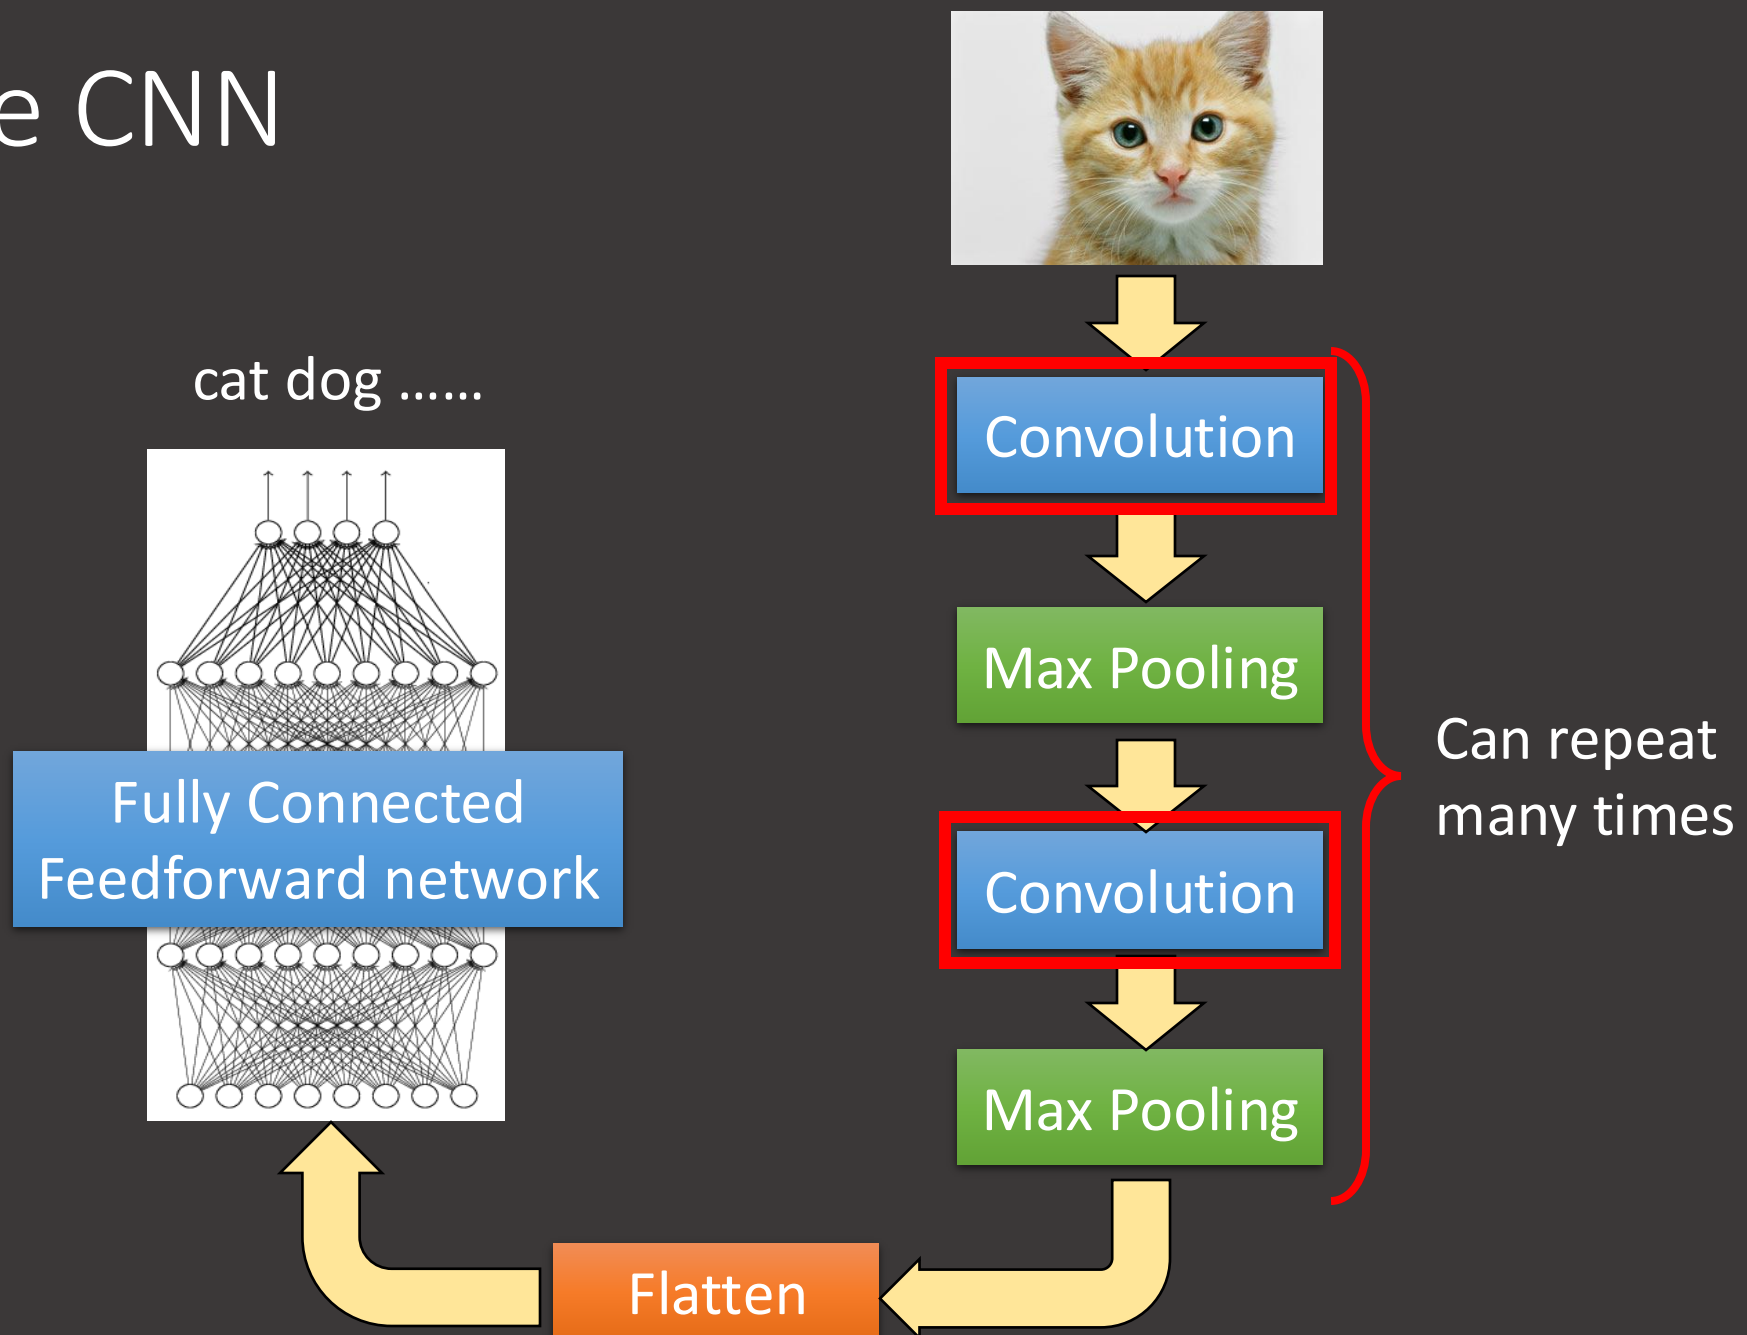
The whole CNN



The whole CNN



The whole CNN



CNN – Convolution

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

Those are the network parameters to be learned.

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1
Matrix

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2
Matrix

⋮

Property 1

Each filter detects a small pattern (3 x 3).

CNN – Convolution


stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1



CNN – Convolution

If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1



We set stride=1 below

CNN – Convolution

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

Property 2

Slide adopted from

<https://www.youtube.com/watch?v=FrKWiRv254g>

CNN – Convolution

stride=1


1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

Do the same process for every filter

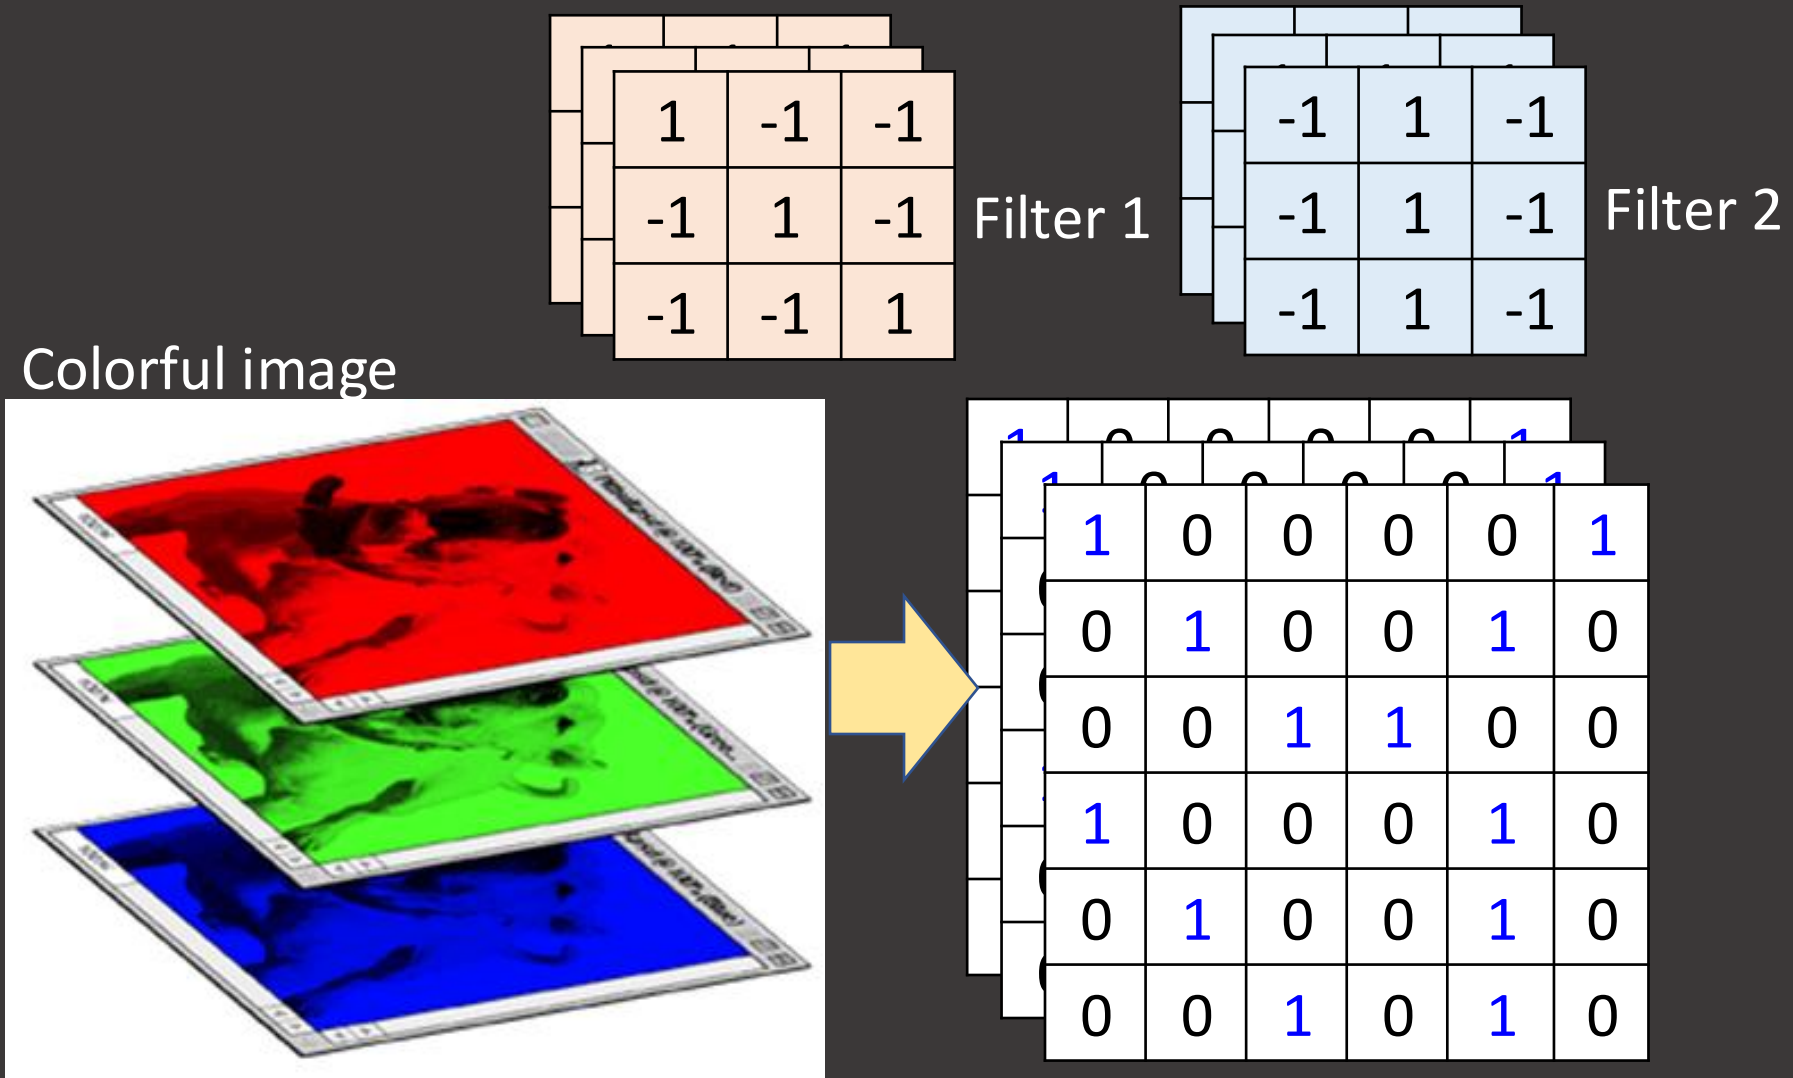


4 x 4 image

Slide adopted from

<https://www.youtube.com/watch?v=FrKWiRv254g>


CNN – Colorful image



Slide adopted from


<https://www.youtube.com/watch?v=FrKWiRv254g>

Convolution v.s. Fully Connected




Fully-
connected

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

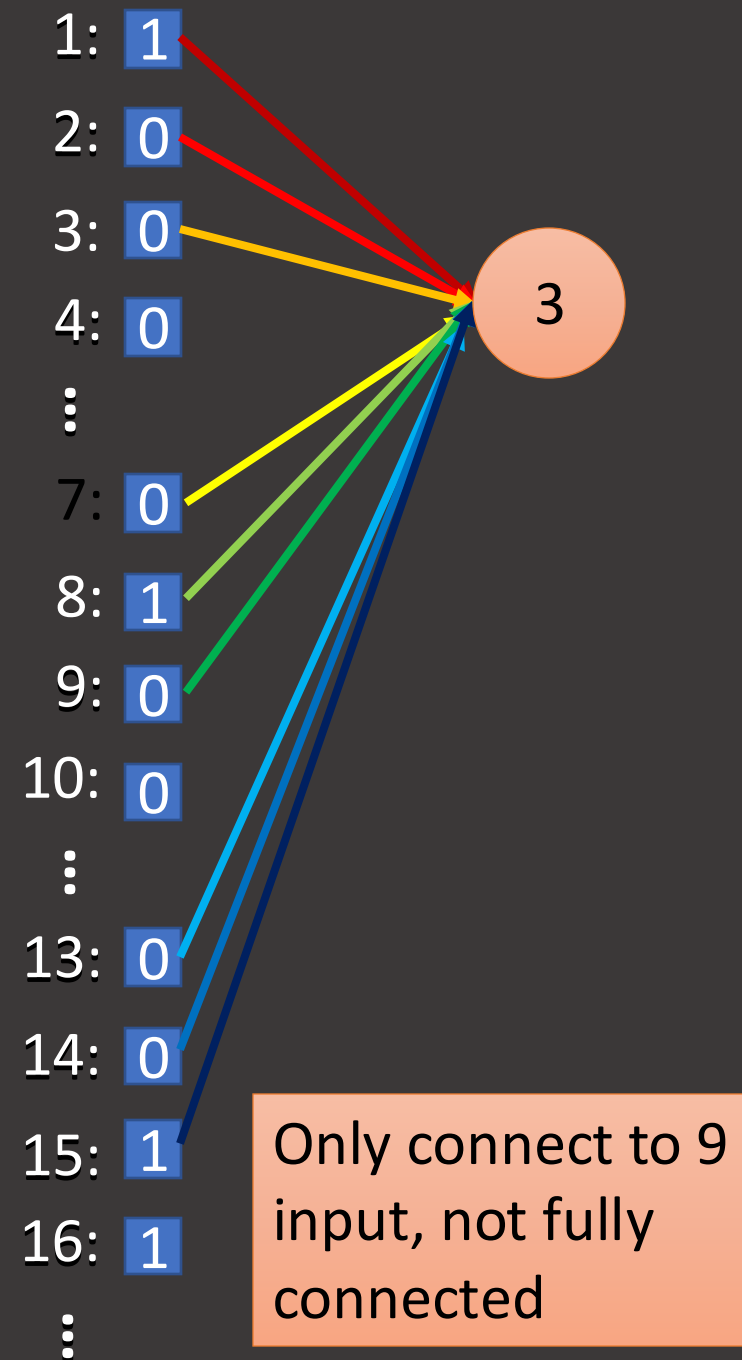


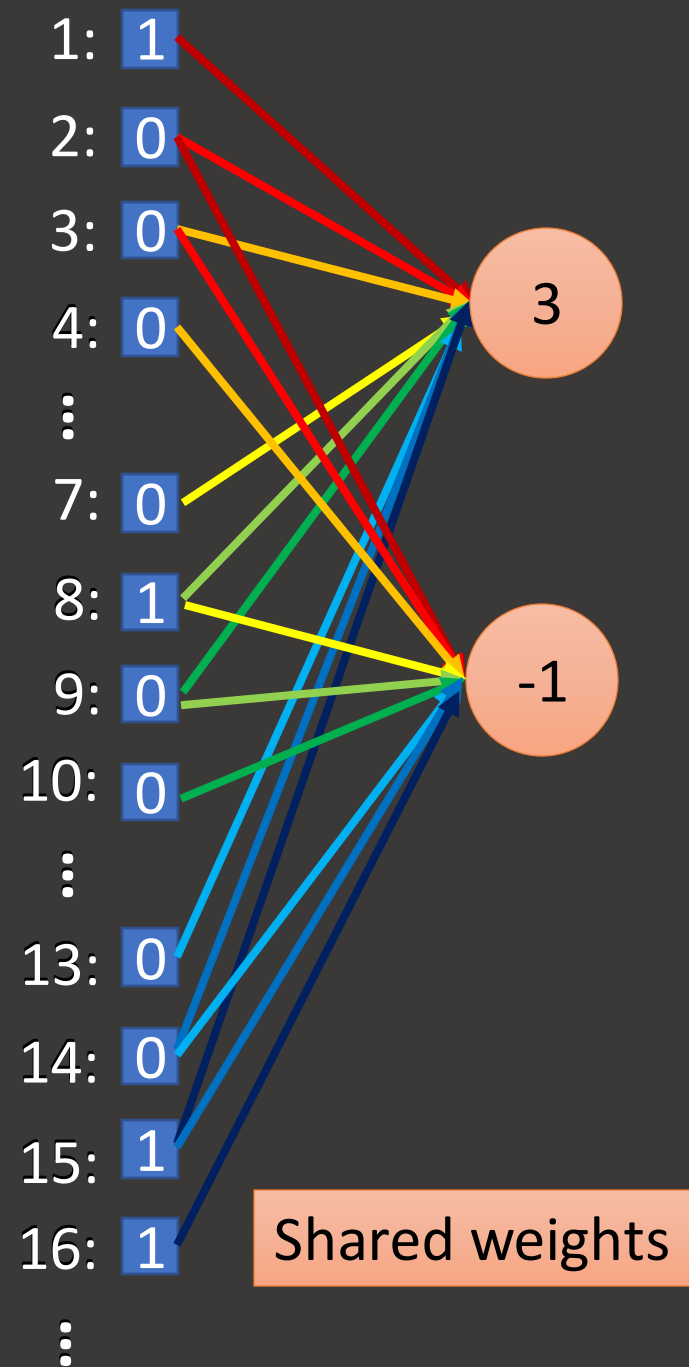

Slide adopted from

<https://www.youtube.com/watch?v=FrKWiRv254g>



Less parameters!





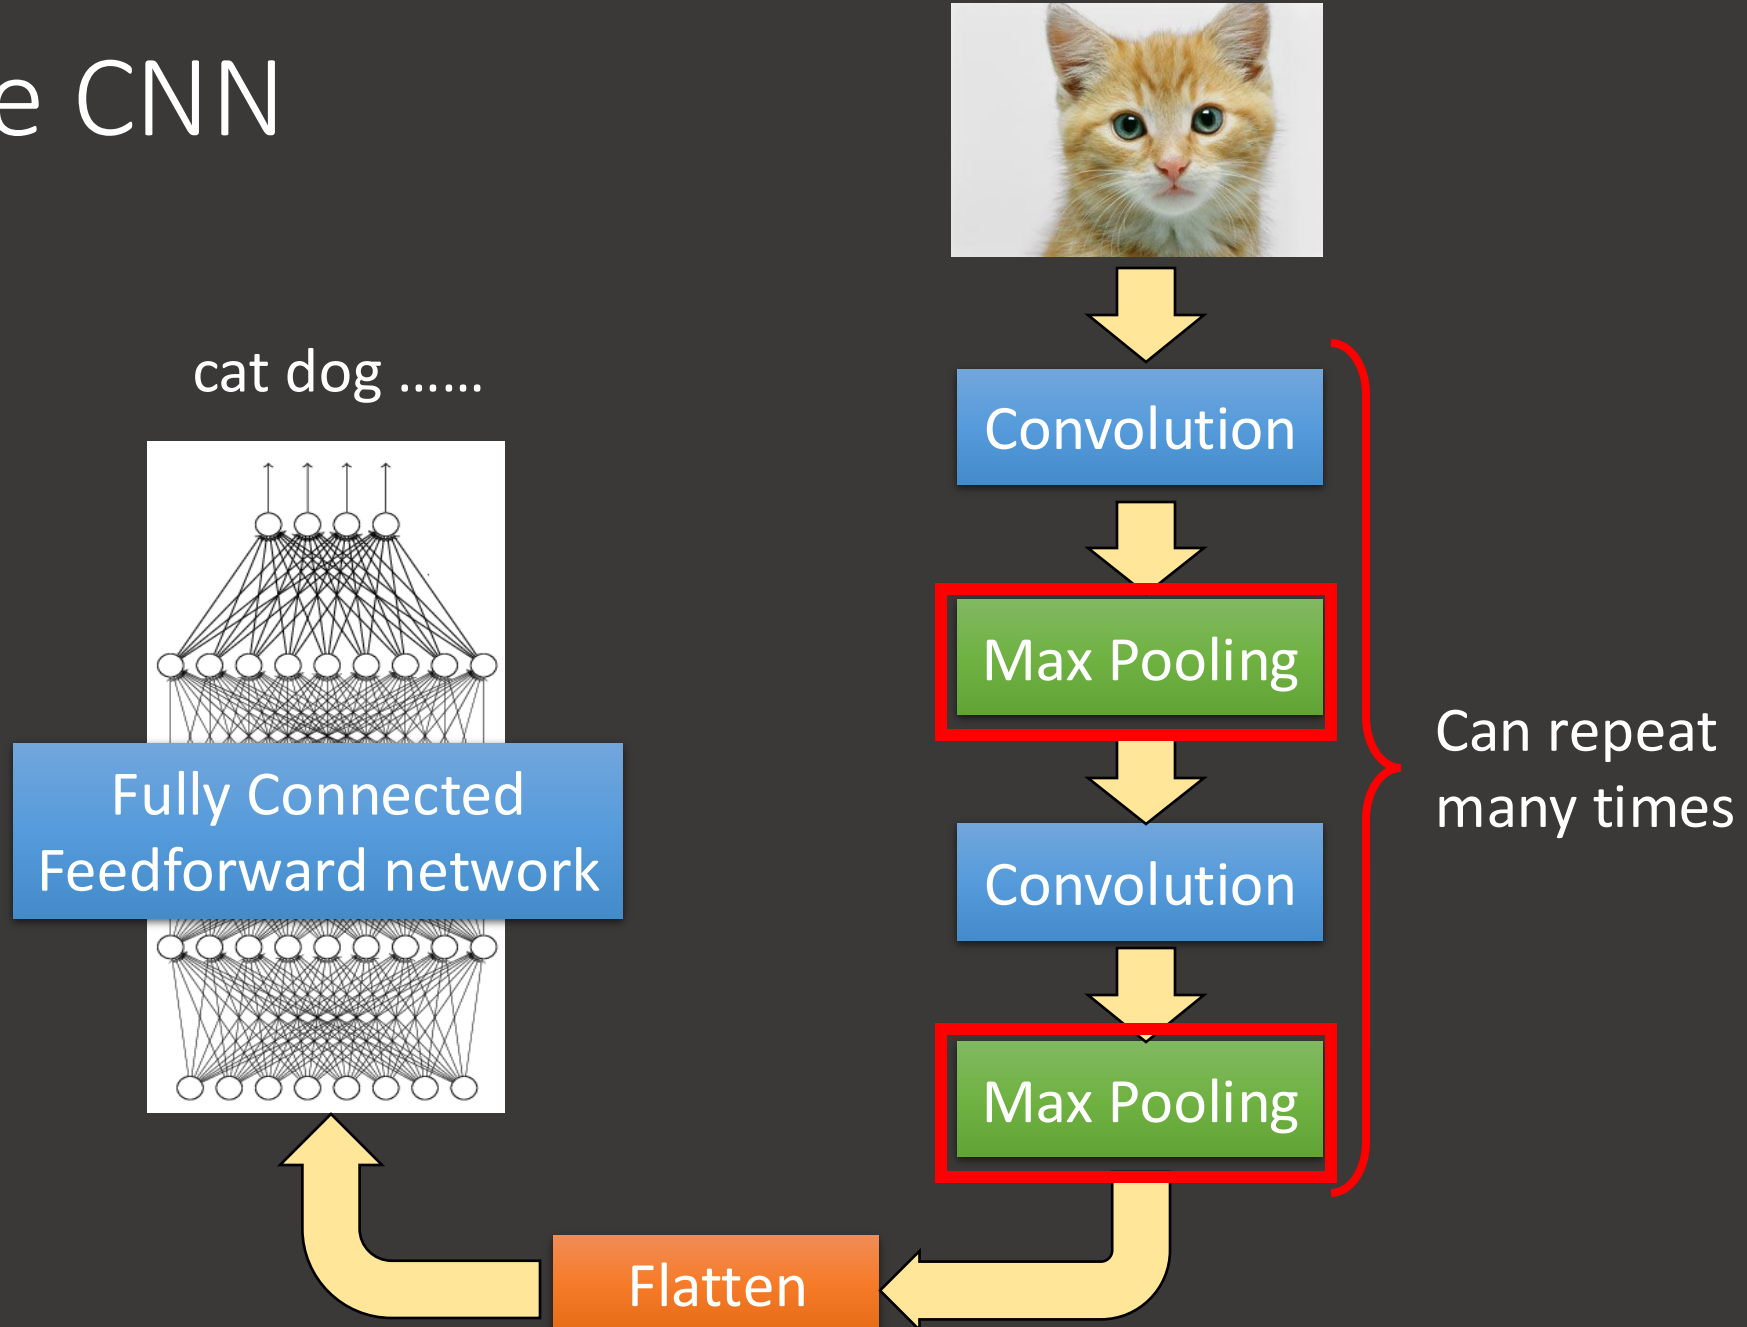
Less parameters!

Even less parameters!

Slide adopted from

<https://www.youtube.com/watch?v=FrKWiRv254g>

The whole CNN



CNN – Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

-1	-1	-1	-1
-1	-1	-2	1
-1	-1	-2	1
-1	0	-4	3


Slide adopted from

<https://www.youtube.com/watch?v=FrKWiRv254g>

CNN – Max Pooling

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0


6 x 6 image




Conv



Max
Pooling



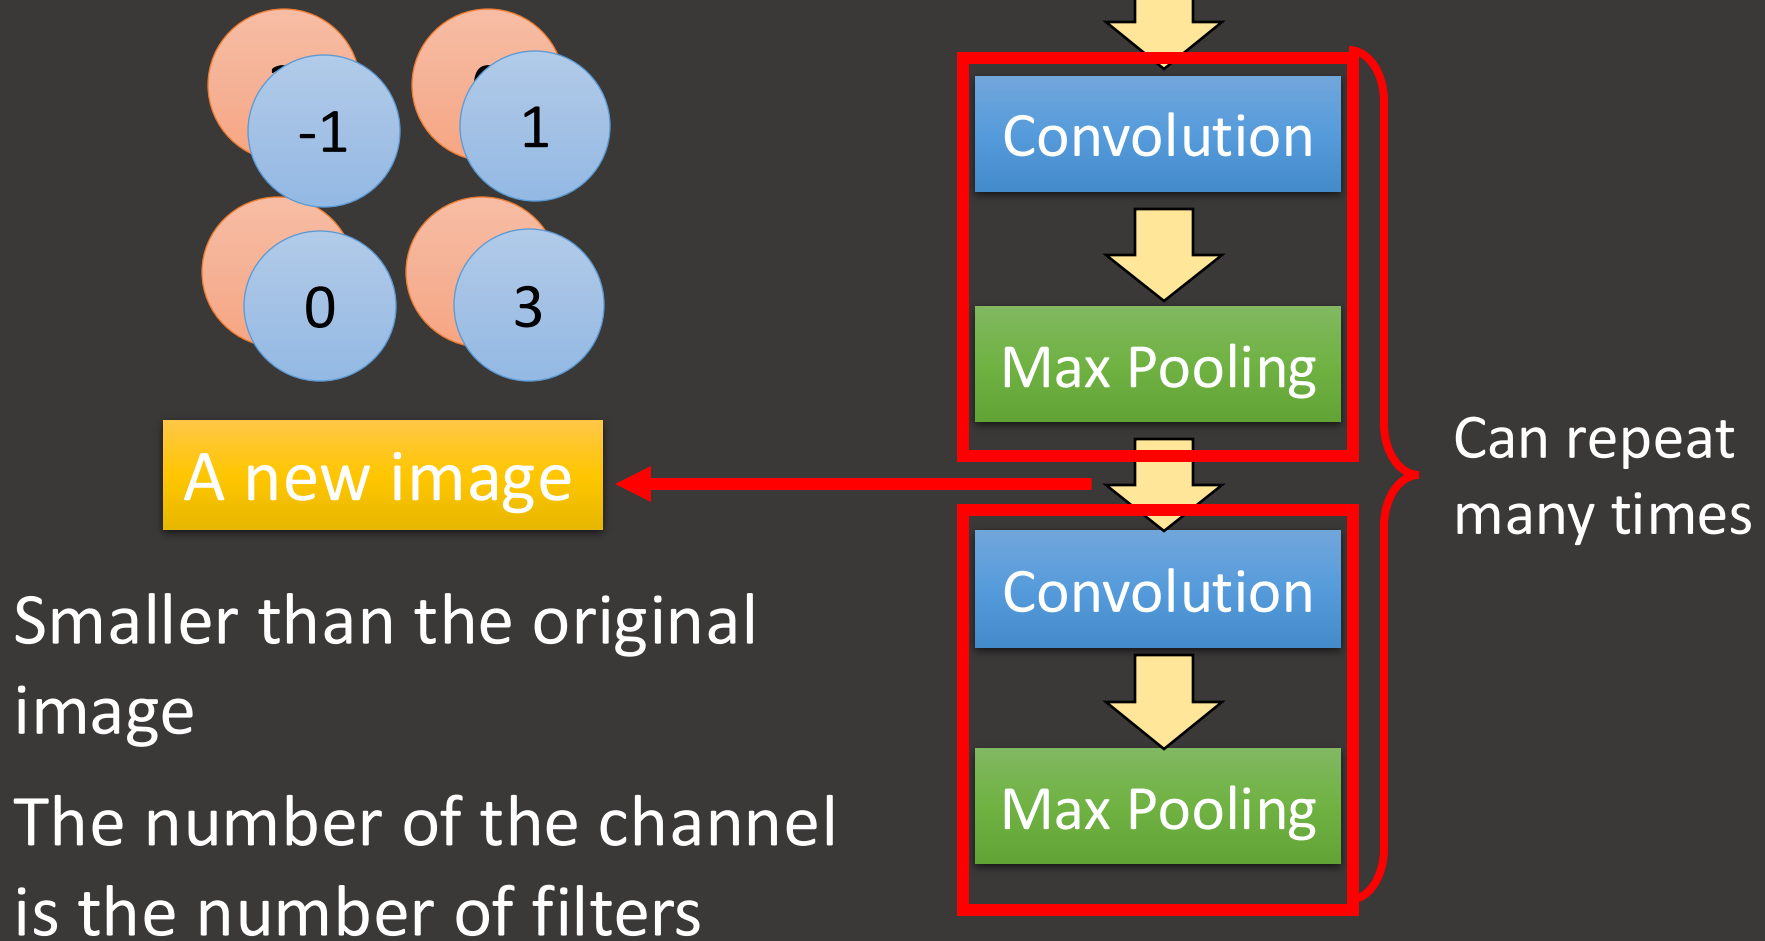
New image
but smaller



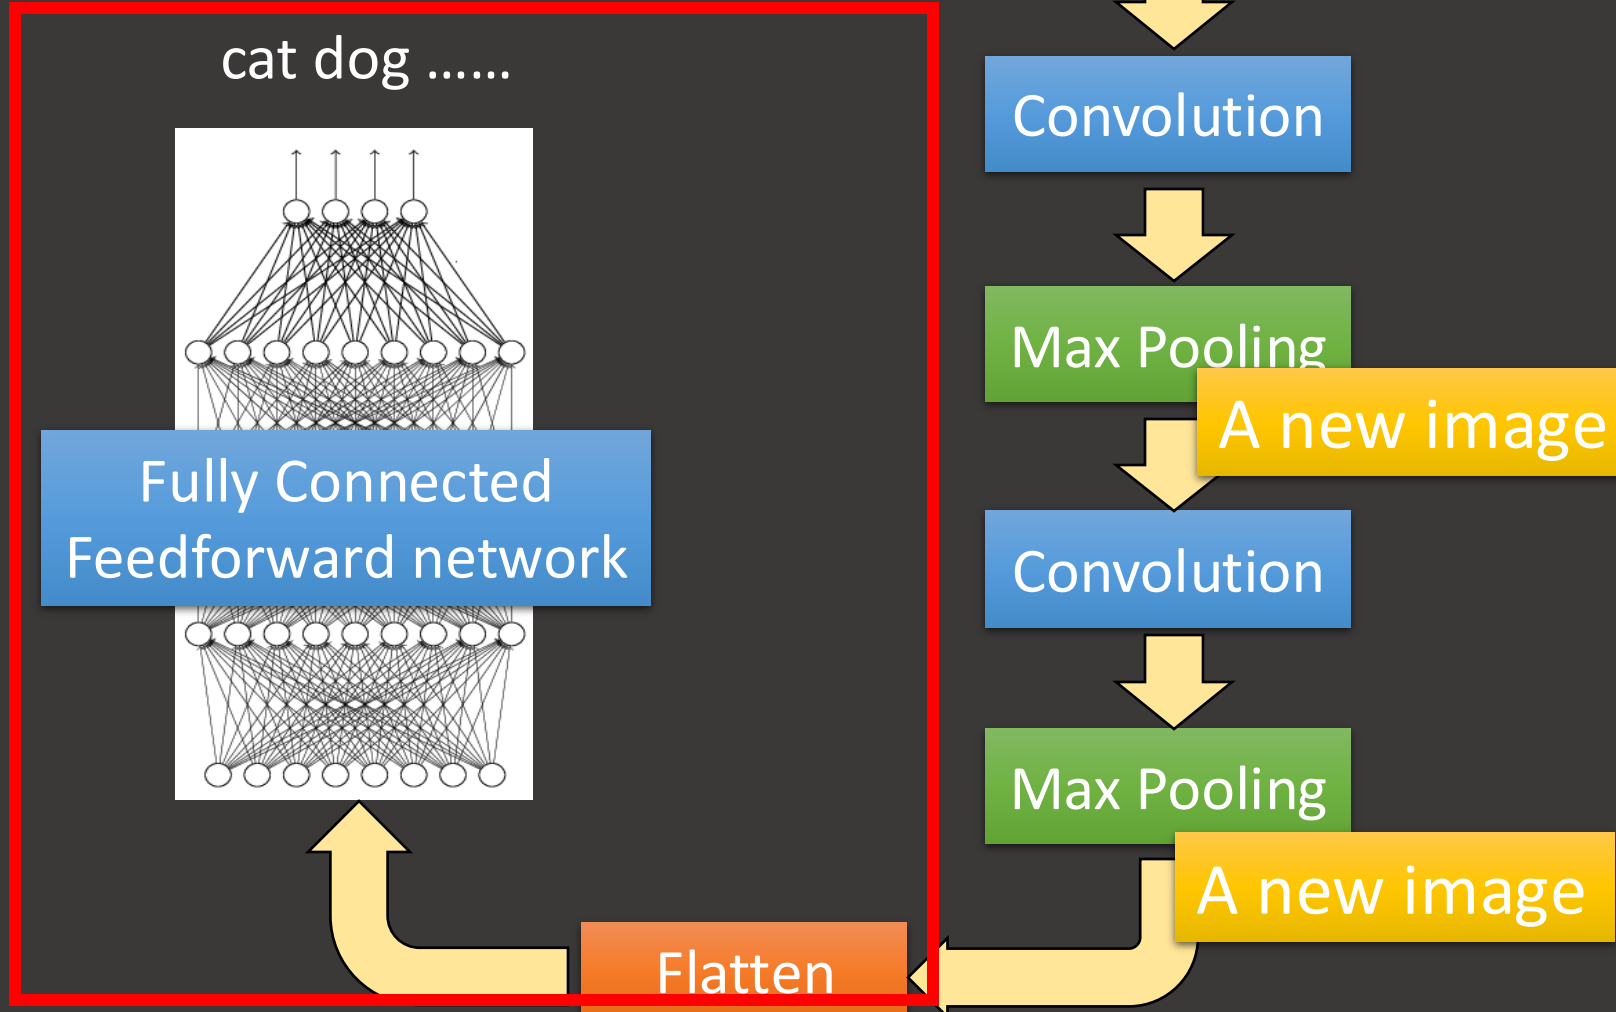
2 x 2 image

Each filter
is a channel

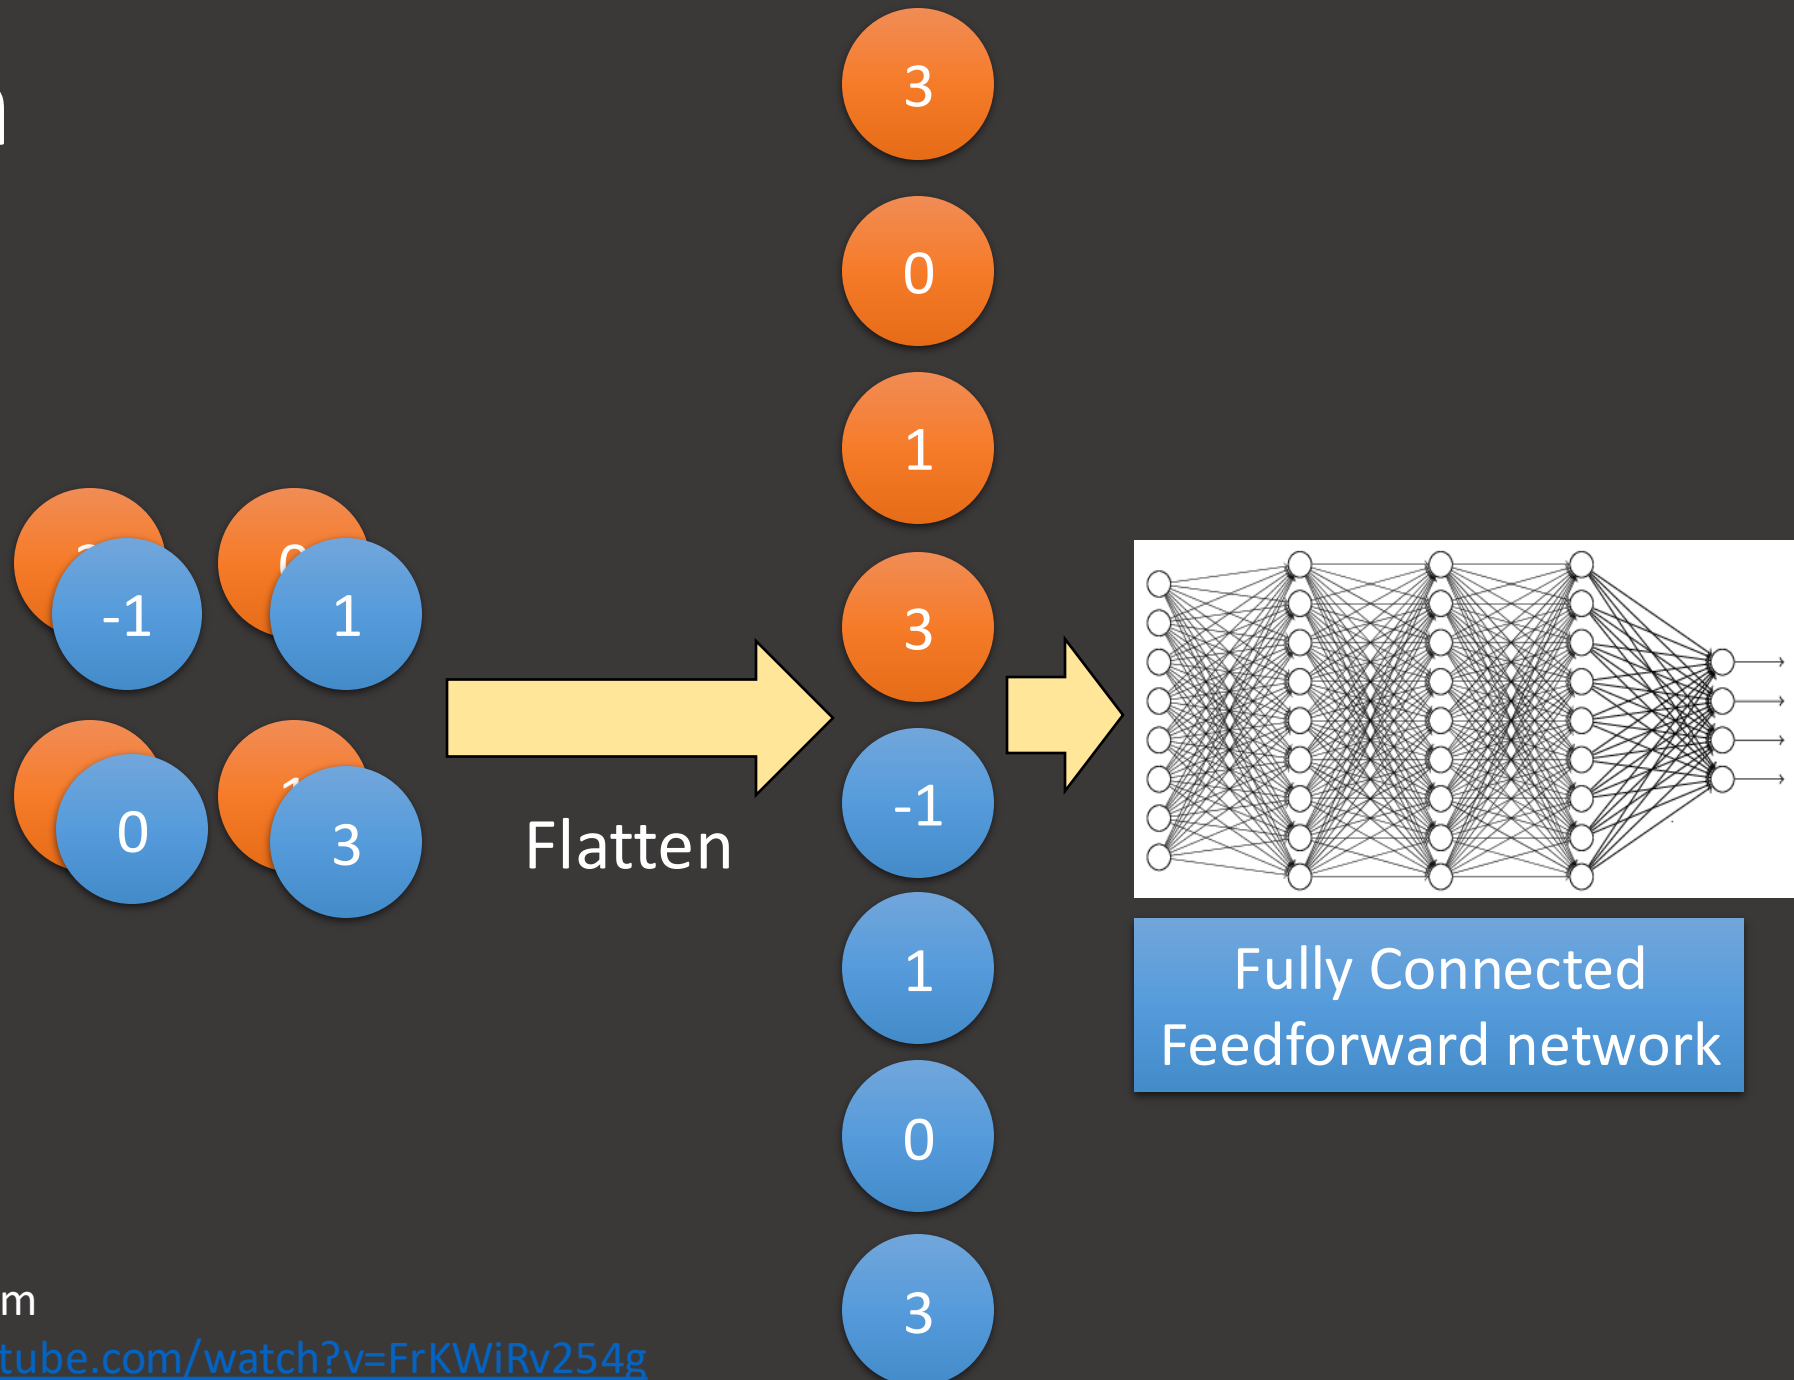
The whole CNN



The whole CNN



Flatten



Slide adopted from


<https://www.youtube.com/watch?v=FrKWiRv254g>

How many layers?

- In general, more the conv layers, greater the extraction of features.
- Example – VGG 16




How many layers?




Why does this happen?

- Increasing the number of layers increases the non-linearity
- Therefore, the problem space becomes increasingly non-convex and thus more difficult to optimize




VGG-16 loss function




VGG-56 loss function

(Drawn using [Loss Landscape Visualizer](#))


ResNet




...



ResNet




Without skip connections




With skip connections

ResNet

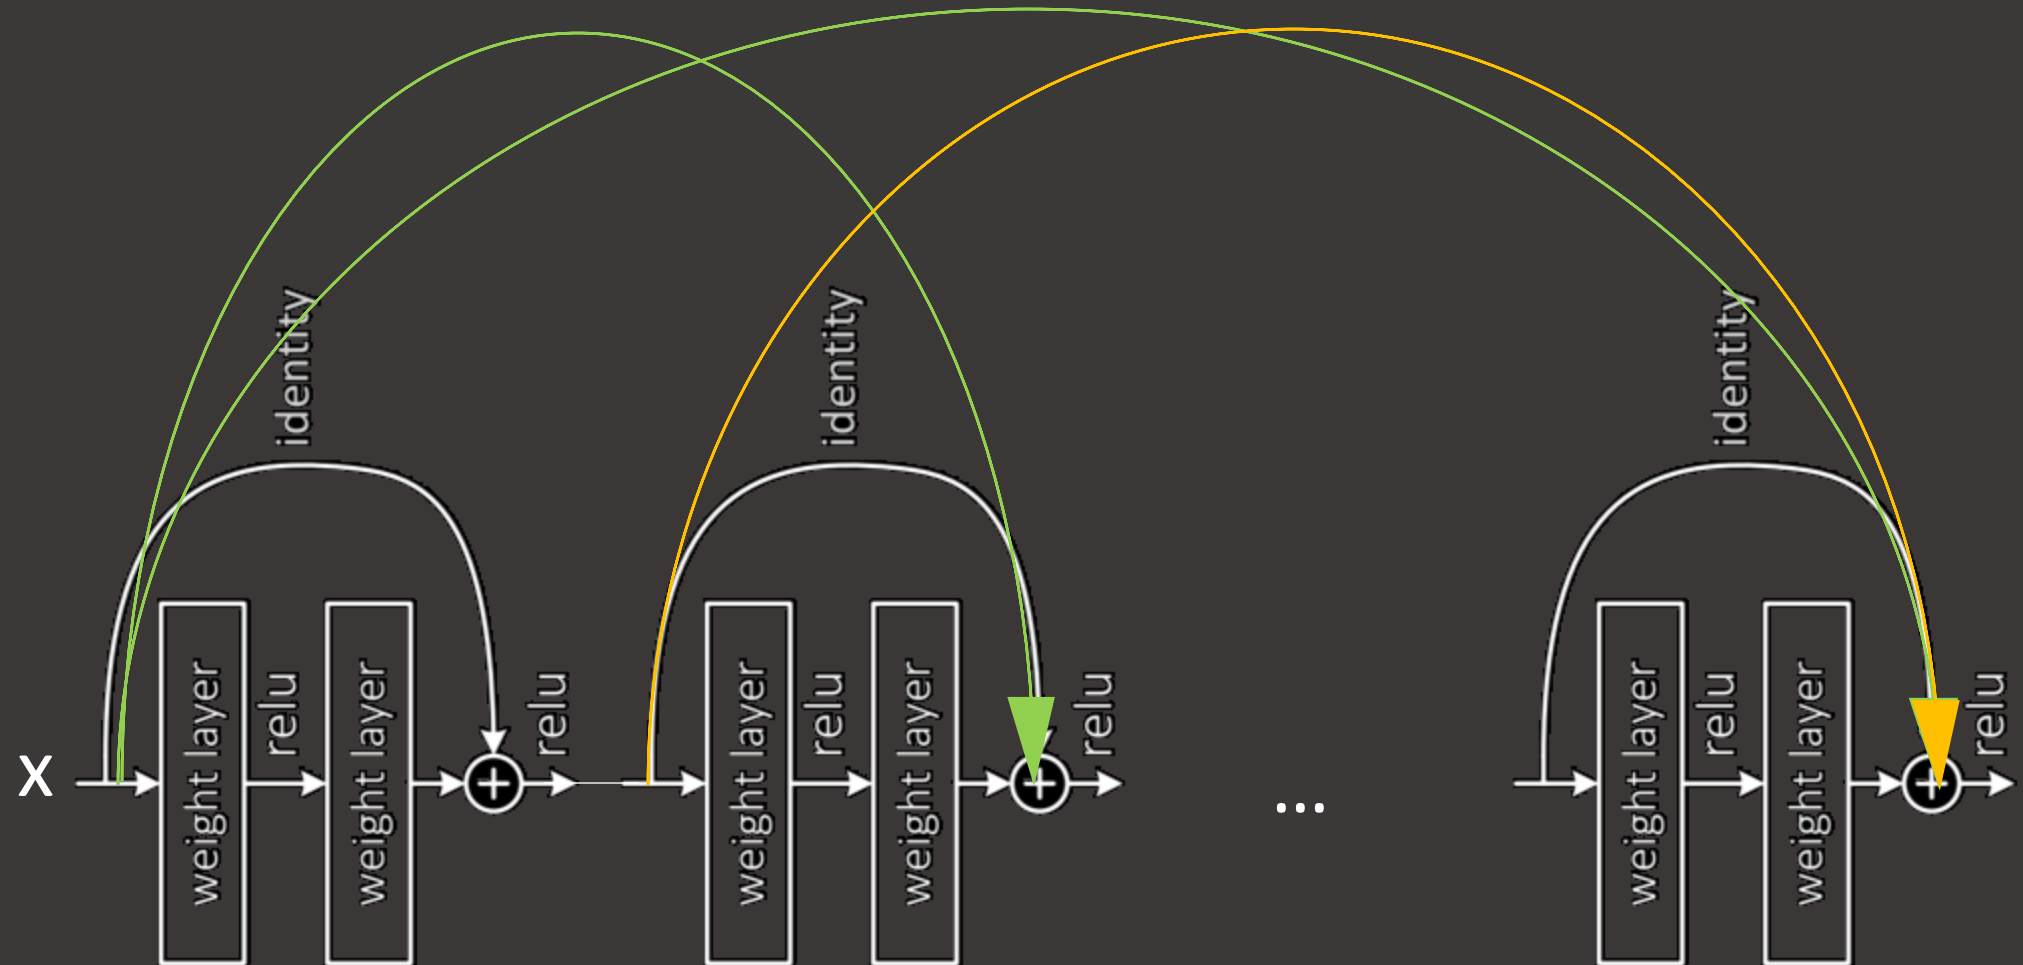
“Memory” gets lost
after one block!!!



...




DenseNet



Huang G, Liu Z, Van Der Maaten L, Weinberger KQ. Densely connected convolutional networks. In Proceedings of the IEEE conference on computer vision and pattern recognition 2017 (pp. 4700-4708).


DenseNet



Deep Learning

- ResNet and DenseNet could be hundreds or thousands of layers
- But that's just the beginning of deep learning

What we have learned so far



Course recommendation: Intro to Reinforcement Learning

Online

Python-based

Dedicated 100-min Q&A Session (Office Hour) per week. Thus, not quite different from in-person course.

Quiz, 3-4 assignments, One group project (2~3 people)

Similar logistics like my other courses.

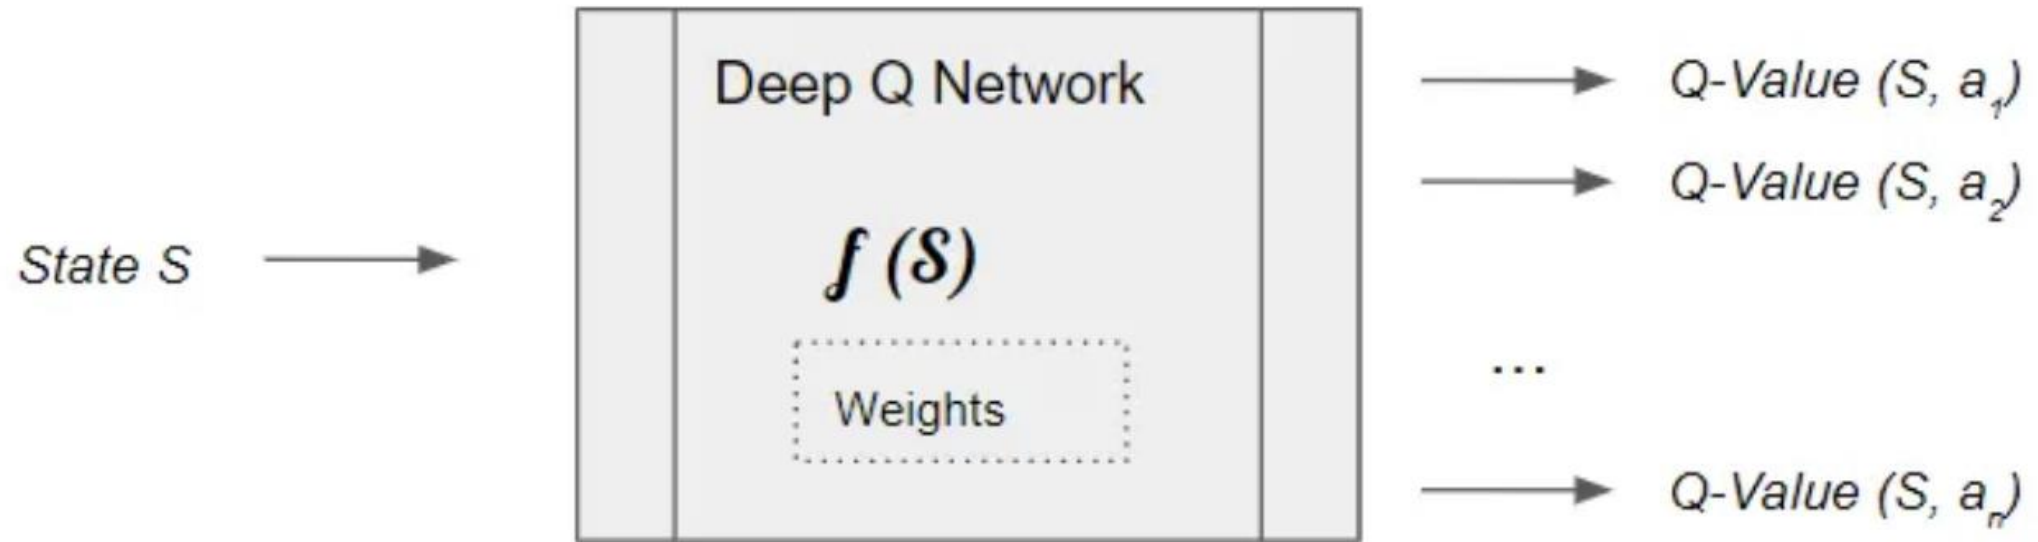
Just need ML pre-requisite, we just remove the DL pre-requisite (need some time to take effect)

Deep RL: RoboCup



<https://youtu.be/xkoXeF9oVH4>

Deep RL: RoboCup



<https://towardsdatascience.com/reinforcement-learning-explained-visually-part-5-deep-q-networks-step-by-step-5a5317197f4b>